

Rishi D. Jha

MASTER'S STUDENT · SECURITY RESEARCHER

☎ 425.677.4846 | ✉ rjha01@cs.uw.edu | 🏠 rishijha.com | 📱 rjha18 | 🌐 rishi-jha

"Master's student, research assistant, and former computer science and mathematics double major. Interests include Machine Learning Fundamentals, Robustness, Security / Privacy, Anomaly Detection, and Graph Theory. Motivated by hard problems."

Education

University of Washington — Seattle

Seattle, WA

MS., COMPUTER SCIENCE

Sep. 2022 - Present

- GPA: In Progress
- Graduate Research Assistant with Dr. Sewoong Oh

University of Washington — Seattle

Seattle, WA

BS.BA., COMPUTER SCIENCE AND MATHEMATICS — PHILOSOPHY: *Cum Laude, Phi Beta Kappa*

Sep. 2018 - Mar. 2022

- GPA: 3.84 / 4.0
- Jun. 2022: Graduated Cum Laude with Phi Beta Kappa honors
- 2018-22: Dean's List (all eligible quarters)

Selected Coursework

Current and Future PhD-Level

Computer Science

Cryptography, Reinforcement Learning, Deep Learning Theory, Randomized Algorithms

Machine Learning (PhD), Cryptography, Deep Learning, NLP, Algorithms, Databases

Mathematics

Real Analysis I & II, Probability and Statistics I, II, & III, Modern Algebra I & II, Linear Algebra

Philosophy

Neuroethics

Publications

WORKSHOP PAPERS

- [1] Dimitrios C. Gklezakos, Rishi Jha, and Rajesh P.N. Rao. "Hyper-Universal Policy Approximation: Learning to Generate Actions from a Single Image using Hypernets". In: *Neurovision 2022: A CVPR Workshop*. New Orleans, USA: Conference on Computer Vision and Pattern Recognition, June 2022. URL: <https://drive.google.com/file/d/113sfBzI69Do9YUJk6oN0j4ycUEaTBHe/view>.
- [2] Rishi Jha and Kai Mihata. "On Geodesic Distances and Contextual Embedding Compression for Text Classification". In: *Proceedings of the Fifteenth Workshop on Graph-Based Methods for Natural Language Processing (TextGraphs-15)*. Mexico City, Mexico: Association for Computational Linguistics, June 2021, pp. 144–149. URL: <https://www.aclweb.org/anthology/2021.textgraphs-1.15>.

PATENTS (PENDING)

- [3] Nisha S. Hameed, Rishi D. Jha, and Evan Argyle. "Graph-Based Analysis of Security Incidents". U.S. pat. Microsoft.

Academic Research

Sewoong Lab — Theoretical Machine Learning and Security

Seattle, WA

GRADUATE RESEARCH ASSISTANT

May 2021 – Present

Working with **Dr. Sewoong Oh** and **Jonathan Hayase** to:

- Develop a novel Neural Tangent Kernel (NTK)-based backdoor attack that persists through the knowledge distillation process and infects networks with triggers they have never seen. The attack only alters the labels for holdout images and preserves good clean test accuracy while infecting a distilled network with a backdoor. Planned submission in **Winter 2023**.
- (Previously) Create an open-source 'backdoor'-attack-benchmark platform and survey for robust machine learning algorithms. Code can be found **here**.

Center for Neurotechnology

Seattle, WA

UNDERGRADUATE ML RESEARCHER

Mar. 2020 – Aug. 2022

Paper accepted at **NeuroVision '22 at CVPR** [1]. Worked with **Dr. Rajesh Rao** and **Dr. Dimitrios Gklezakos** to:

- Develop a low-cost, 'personalized' hypernetwork for hierarchical and task-conditional RL called the Hyper-Universal Policy Approximator (HUPA). HUPAs are up to 35% more resilient to sparsity and have up to 25% generalization than their traditional embedding alternatives. Planned full conference submission in **Winter 2023**.
- Construct an audio-visual hypernetwork for representation learning and classification on a massive dataset in which a video-controlled neural network controls the weights of an audio interpreter.
- Create a convolutional, manifold-learning based network to learn complex features in natural images in an unsupervised fashion using sparse coding. The system learns representational similarities between features and generalizes them.

Self-Directed

Seattle, WA

NLP RESEARCHER

Nov. 2020 – Jun. 2021

Paper accepted at **TextGraphs '21 at NAACL** [2]. Worked with **Kai Mihata** to:

- Investigate the downstream effects of compressing BERT embeddings using nonlinear dimensionality reduction techniques and geodesic estimations.
- Find that nonlinear compressions of the embeddings tend to work well in some data regimes, a feature that can be utilized in memory-constrained settings.

ICTD Lab

Seattle, WA

UNDERGRADUATE RESEARCHER

Nov. 2018 - May 2019

Worked with **Dr. Spencer Sevilla** to:

- Investigate the performance dynamics of different chat apps in poor network conditions.
- Implement a teaching solution for schoolchildren in rural Indonesia.

Research in Industry

Microsoft Defender Research

Redmond, WA

SOFTWARE ENGINEERING INTERN — DATA SCIENCE

Jul. 2022 - Sep. 2022

- Ideated, pitched, and implemented a low-cost, humanly interpretable meta-learning framework that exploits spectral similarities in existing classifier responses to drive robustness in the Defender product. The productionalized system was lightweight, had upwards of 97% precision and recall, and was humanly interpretable.
- The model is being pushed from pre-production to production and will start providing protection for billions of users by **Summer 2023**.

Microsoft Defender Research

Remote

SOFTWARE ENGINEERING INTERN — DATA SCIENCE

Jun. 2021 - Sep. 2021

Patent submitted in **Winter 2022** [3].

- Ideated and designed patent-pending approach to detect malicious Command-and-Control intrusions in corporate networks using spectral methods on graphs. The model achieved high precision and recall in finding Indicators of Compromise in historical data.
- The project has received significant investment from the team since my departure with a goal of pushing an extension of the model to production in **Summer 2023**.

Teaching

University of Washington — Seattle

Seattle, WA

4X UNDERGRAD / GRAD MACHINE LEARNING TA

Mar. 2020 - Dec. 2021

During Spring 2020, Winter 2021, Spring 2021, Autumn 2021:

- Taught undergraduate and graduate students as an undergraduate through 25-person sections and biweekly office hours.
- Designed section materials for entire teaching staff, monitored discussion boards, and graded assignments.

University of Washington — Seattle

Seattle, WA

MACHINE LEARNING COURSE DESIGNER

Jun. 2021 - Sep. 2021

During Summer 2021, funded by **Dr. Sewoong Oh** to:

- Redesign the course's problem sets and homework infrastructure to keep up with a rapidly evolving course and field.
- Write new problems and solutions for future quarters.
- Create a new central grading system and TA codebase for future quarters and course staffs to use.

Other Work Experience

Microsoft

Remote

SOFTWARE ENGINEERING INTERN — DEFENDER SECURITY

Jun. 2020 - Sep. 2020

- Reduced related COGS by \$100K - \$1M by creating ML model to selectively download dangerous files for analysis. In production.
- Built infrastructure for safer ML model deployment. In production.
- Decreased researcher rule development time by 35%, by creating VSCode extension to natively test rules. In production.

Microsoft

Redmond, WA

EXPLORE INTERN — OFFICE.COM FRONT END

Jun. 2019 - Aug. 2019

- Designed, implemented, and released front end notes tool for the Office.com team using Typescript, Redux, and React internally.

Honors

2022	Appointed , Phi Beta Kappa	Seattle, WA
2022	Appointed , Cum Laude Scholar	Seattle, WA
2018-22	Selected , Dean's List (all eligible quarters)	Seattle, WA
2021-22	Selected , Varsity Climbing Team at UW	Seattle, WA
2020	1st Place , Rain City Send Bouldering Competition — Recreational Category	Seattle, WA
2019	Finalist , (Top 4 of 36 Teams) UW Foster CBDC: Consulting Challenge	Seattle, WA
2017	3rd Place , (1000+ Teams) Microsoft OneWeek Hackathon Consumer Category	Redmond, WA

Skills

Interests	Machine Learning, Robustness, Security, Privacy, Anomaly Detection, Graph Theory
Technical	Python, PyTorch, TensorFlow, JAX, C++, Java / C#,
Languages	English, Hindi, Spanish